



MongoDB

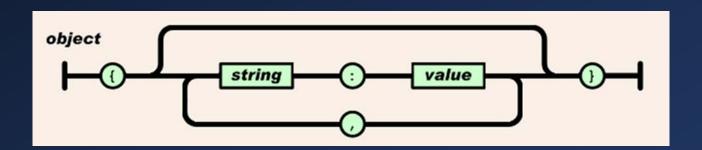
- A document-oriented database, also known as document store database, aggregate database, or simply document store or document database
- A database that uses a <u>document-oriented model</u> to store data
- Document—oriented databases store each record and its associated data within a single document.
- Each document contains semi-structured data that can be queried against using various query and analytics tools
- MongoDB stores data records as <u>BSON documents</u>. BSON is a binary representation of <u>JSON</u> documents.

JSON

- JSON (JavaScript Object Notation, http://www.json.org/) is a lightweight data-interchange format.
- JSON is text, written with JavaScript object notation.
- JSON is built on two structures:
 - A collection of name/value pairs.
 - In various languages, this is realized as an object, record, structure, dictionary, hash table, keyed list, or associative array.
 - An ordered list of values.
 - In most languages, this is realized as an array, vector, list, or sequence.
- Virtually all modern programming languages support them in one form or another.

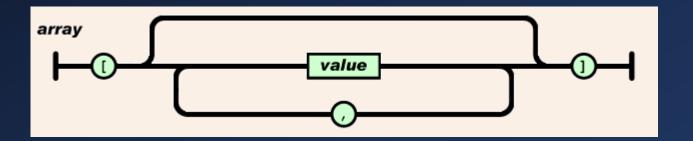
JSON - object

- An object is an unordered set of name/value pairs.
 - An object begins with { (left brace) and ends with } (right brace). Each name is followed by: (colon) and the name/value pairs are separated by, (comma).



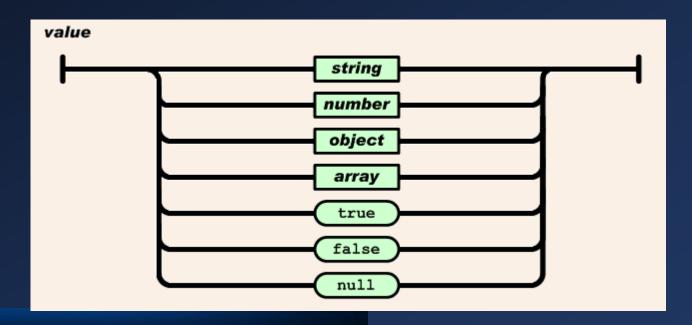
JSON - array

- An array is an ordered collection of values.
 - An array begins with [(left bracket) and ends with] (right bracket). Values are separated by , (comma).



JSON - value

 A value can be a string in double quotes, or a number, or true or false or null, or an object or an array. These structures can be nested.



JSON example

- How many objects are there in this JSON example?
 - 5
 - 6 ✓
 - 8
 - 11
 - Other answers

MongoDB Document Structure

 MongoDB documents are composed of field-and-value pairs and have the following structure:

```
field1: value1,
  field2: value2,
  field3: value3,
   ...
  fieldN: valueN
}
```

```
{
    _id: ObjectId("5099803df3f4948bd2f98391"),
    name: { first: "Alan", last: "Turing" },
    birth: new Date('Jun 23, 1912'),
    death: new Date('Jun 07, 1954'),
    contribs: [ "Turing machine", "Turing test", "Turingery" ],
    views : NumberLong(1250000)
}
```

 The value of a field can be any of the BSON data types, including other documents, arrays, and arrays of documents.

Data type	Description
String	text values, e.g. "Glasgow"
Boolean	TRUE or FALSE.
Integer (32b and 64b):	numerical value, e.g. 5
Double	floating point values, e.g. 5.24
Min / Max keys	values which will always compare lower / higher than any other type
Timestamp	used to store a timestamp, 64-bit integer
Null	used for a Null value
Date	used to store the current date or time in UNIX time format (POSIX time), 64-bit integer
Object ID	used to store the document's ID, generated by database
Binary data	used to store binary data.
Regular expression	
JavaScript Code	

MongoDB Document Structure

- Field names are strings.
- Restrictions on field names:
 - The field name _id is reserved for use as a primary key; its value must be unique in the collection, is immutable, and may be of any type other than an array.
 - The field names **cannot** start with the dollar sign (\$) character.
 - The field names **cannot** contain the dot (.) character.
 - The field names cannot contain the null character.
- Field Value Limit

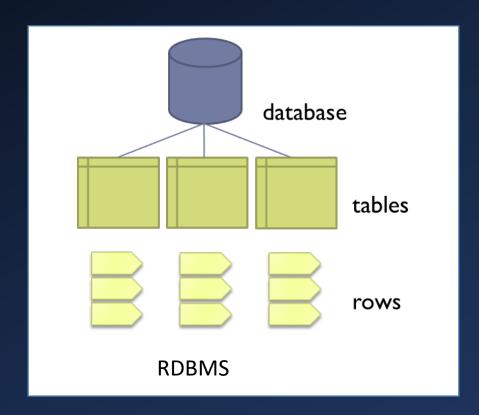
For indexed collections, the values for the indexed fields have a Maximum Index Key Length limit.

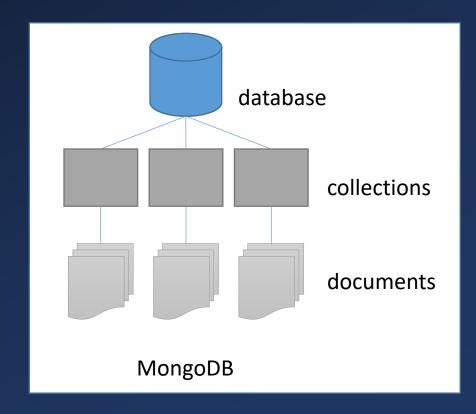
MongoDB vs Relational Database

Terminology

RDBMS	MongoDB			
Database	Database			
Table	Collection			
Tuple/Row	Document			
column	Field			
Table Join	Embedded Documents			
Primary Key	Primary Key (Default key _id provided by mongodb itself)			
Database Server and Client				
Mysqld/Oracle	mongod			
mysql/sqlplus	mongo			

Model





RDMBS

- Meaning of data items in each row is defined by the design of the table that contains the row
- Column names, types, constraints
- Database has a well defined structure, or schema, for the data it can store
- Database will reject data that doesn't conform to schema

MongoDB

- Meaning of data items in a document is contained in the document itself – name part of name/value pairs
- Database/collections have no knowledge of meaning of data
- Database (like many NoSQL stores) is schemaless
- Database will allow documents with any structure to be stored
- Can store unstructured data where the content is not known apriori and can vary from document to document

Getting start with MongoDB

- Installation
 - Download from:

https://www.mongodb.com/download-center?initial=true#community

• Install manual:

https://docs.mongodb.com/manual/administration/install-community/

- Tools in bin folder after installation:
 - |-- bin
 - | |-- mongo (the database shell)
 - | |-- mongod (the core database server)
 - | |-- mongos (auto-sharding process)
 - | |-- mongodump (dump/export utility)
 - | |-- mongorestore (restore/import utility)

 MongoDB requires a data folder to store its files. MongoDB will not start if a data directory doesn't exist.

- Default data directory is:
 - /data/db (Linux, MacOS)
 - c:\data\db (Windows)

May need to set permissions to allow user to read & write

- Start server with command at system command prompt: mongod --dbpath data/db
 - --dbpath option allows you to specify data directory, can omit if using the default location, example here assumes data directory is <u>data/db</u> in the current working directory
 - Assuming here that mongodb bin folder is in path, may need to specify full path otherwise
- You should see startup messages including something like this, tells you the server is alive and what TCP port number is being used (27017)

2018-09-31T14:36:40.701+0000 | NETWORK [thread1] waiting for connections on port 27017

- Note that you now need to keep the terminal window open, server will shut down if you close it
- You can also install and run MongoDB as a service, detached from any terminal window

 From another terminal on the same computer, enter the following command to run mongo shell client:

mongo

- This assumes that there is no security configured in the database, which is OK for learning. But in a "real" installation, username/password would need to specify see documentation
- Will look for server instance running on localhost at default port (27017)
- Should see the following message and the > shell prompt (may see other messages also):

```
MongoDB shell version v3.4.0 connecting to: mongodb://127.0.0.1:27017 MongoDB server version: 3.4.0 >
```

 MongoDB server can also be accessed from a remote machine with mongo client installed, e.g.

```
mongo --host 192.168.231.127 --port 27017
```



MongoDB CRUD Operations

Creating databases/collections using the shell

```
db shows the name of the current database use myDatabase switches to a different database, creates if it doesn't exist db.myCollection.insert( {"country":"Scotland"} ) inserts a document in the specified collection in the current database, creates collection if it doesn't exist
```

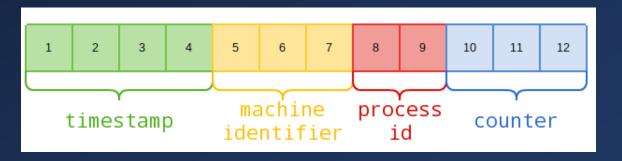
• Since version 3.2, MongoDB provides the following methods to insert documents into a collection:

```
db.myCollection.insertOne()
db.myCollection.insertMany()
```

You can check that the document has been stored using the find function:

```
> db.myCollection.find()
{ "_id" : ObjectId("59e01dd19400bbe099b3b35d"), "country" : "Scotland" }
{ "_id" : ObjectId("59e01ea59400bbe099b3b35e"), "country" : "England",
    "city" : "London" }
{ "_id" : ObjectId("59e0200d9400bbe099b3b35f"), "city" : "Washington",
    "sport" : "basketball", "country" : "USA"}
```

- We didn't have to specify what properties a document can have, and that different documents can have different fields.
- MongoDB adds an _id property to each document, if you do not specify the _id value manually, then the type will be set to a special BSON datatype that consists of a 12-byte binary value:



Update an existing document by adding a new property:

```
db.myCollection.updateOne(
    {"country" : "USA"},
    {$set : {"city" : "New York"}})
```

Delete a specific document:

```
db.myCollection.remove({"city" : "London"})
```

Deleting a whole collection from the database:

```
db.myCollection.drop()
```

Data Retrieval in RDBMS

- Query language
 - A language in which a retrieval request is specified, i.e., to specify the data items you need.
- SQL Structured Query Language
 - The ubiquitous query language when the data is structured
- Basic structure of an SQL query:

SELECT (output attribute(s))
FROM (table(s) to use)
WHERE (the condition(s) to satisfy)

Data Retrieval in MongoDB

- MongoDB does not support a query language such as SQL
- Main query mechanism is the find() function
- The basic syntax of find() method is:

```
>db.collection.find()
```

all the documents in the collection will be displayed in a non-structure way.

• The pretty() method:

```
>db.collection.find().pretty()
```

documents will be displayed in a formatted way.

SQL select and MongoDB find()

db.collection.find(<query filter>, ction>).<cursor modifier>

Like FROM clause, specifies the collection to use

Like WHERE clause, specifies which documents to return Projection variables in SELECT clause

How many results to return etc.



The find() Function

- To build a query on a collection, you specify a query filter (or, query document) with properties you wish the results to match and pass this as a parameter to find
- Query document is a JSON document whose properties are either field names of documents in the collection or MongoDB operators operators have names that start with \$, e.g. \$lt, \$exists
- Other functions can be chained with *find* to process the results, e.g. sort, limit, pretty

RDBMS Where Clause Equivalents in MongoDB

Operation	Syntax	Example	RDBMS Equivalent
Equality	{ <key> : <value>}</value></key>	<pre>db.mycol.find({"by":"tutorials point"})</pre>	where by = 'tutorials point'
Less Than	{ <key> : {\$lt:<value>}}</value></key>	db.mycol.find({"likes":{\$lt:50}})	where likes < 50
Less Than Equals	{ <key> : {\$lte:<value>}}</value></key>	db.mycol.find({"likes":{\$lte:50}})	where likes <= 50
Greater Than	{ <key> : {\$gt:<value>}}</value></key>	db.mycol.find({"likes":{\$gt:50}})	where likes > 50
Greater Than Equals	{ <key> : {\$gte:<value>}}</value></key>	db.mycol.find({"likes":{\$gte:50}})	where likes >= 50
Not Equals	{ <key> : {\$ne:<value>}}</value></key>	db.mycol.find({"likes":{\$ne:50}})	where likes != 50

Collection <u>restaurants</u> for query examples

```
" id" : ObjectId("589789858041723d5ee77780"),
               "address" : {
                         "building" : "1007",
                         "coord": [-73.856077,40.848447],
object
                         "street" : "Morris Park Ave",
                         "zipcode" : "10462"
                                                                     object
               "borough": "Bronx",
               "cuisine": "Bakery",
               "grades" : [
                                   "date" : ISODate("2014-03-03T00:00:00Z"),
                                   "grade" : "A",
                                   "score" : 2
       Array
                                   "date" : ISODate("2011-03-10T00:00:00Z"),
                                   "grade" : "B",
                                   "score" : 14
               "name" : "Morris Park Bake Shop",
               "restaurant id" : "30075445"
```

Query on a single field

- Query on multiple fields
 - MongoDB treats each property as having an implicit boolean AND.

>db.restaurants.find({cuisine : "American", borough : "Bronx" }).limit(2)

```
> db.restaurants.find({cuisine : "American", borough : "Bronx" }).limit(2)
{ "_id" : ObjectId("59e565aa8b8edcdd5b7a43c5"), "address" : { "building" : "2300
", "coord" : [ -73.8786113, 40.8502883 ], "street" : "Southern Boulevard", "zipc
ode" : "10460" }, "borough" : "Bronx", "cuisine" : "American", "grades" : [ { "d
ate" : ISODate("2014-05-28T00:00:00Z"), "grade" : "A", "score" : 11 }, { "date"
: ISODate("2013-06-19T00:00:00Z"), "grade" : "A", "score" : 4 }, { "date" : ISOD
ate("2012-06-15T00:00:00Z"), "grade" : "A", "score" : 3 } ], "name" : "Wild Asia
", "restaurant_id" : "40357217" }
{ "_id" : ObjectId("59e565aa8b8edcdd5b7a43f6"), "address" : { "building" : "658"
, "coord" : [ -73.81363999999999, 40.82941100000001 ], "street" : "Clarence Ave"
, "zipcode" : "10465" }, "borough" : "Bronx", "cuisine" : "American", "grades" :
    [ { "date" : ISODate("2014-06-21T00:00:00Z"), "grade" : "A", "score" : 5 }, { "date" : ISODate("2012-07-11T00:00:00Z"), "grade" : "A", "score" : 10 } ], "name"
: "Manhem Club", "restaurant_id" : "40364363" }
```

 MongoDB supports boolean OR queries, but you must use a special operator (\$or) to achieve it.

```
>db.restaurants.find({$or: [{restaurant_id : "30075445"}, {restaurant_id : "30191841"}]})
```

```
db.restaurants.find({$or: [{restaurant_id : "30075445"}, {restaurant_id : "301
        : ObjectId("59e565aa8b8edcdd5b7a43bb"), "address" : { "building"
                -73.856077, 40.848447 l, "street" : "Morris Park
                                     } l, "name" : "Morris Park Bake Shop",
                         "score" :
                   "59e565aa8b8edcdd5b7a43bd"), "address" :
                         "). "grade"
lds Pub And Restaurant", "restaurant id
```

Sgt opertor (greater than) on a field value

>db.movie.find({year: {\$gt:2005}})

```
> db.movie.find({year: {$gt:2005}})
{ "_id" : ObjectId("59e92e4d8eeffc9227fa4a83"), "imdb" : "tt0796366", "title" :
"Star Trek", "year" : 2009, "type" : "movie" }
{ "_id" : ObjectId("59e92e4d8eeffc9227fa4a84"), "imdb" : "tt1408101", "title" :
"Star Trek Into Darkness", "year" : 2013, "type" : "movie" }
{ "_id" : ObjectId("59e94bca8eeffc9227fa4a8b"), "imdb" : "tt0769366", "title" :
"Ice Age 2: The MEltdown", "year" : 2006, "type" : "cartoon" }
{ "_id" : ObjectId("59e94bca8eeffc9227fa4a8c"), "imdb" : "tt1401081", "title" :
"Ice Age 3: Dawn of the Dinosaurs", "year" : 2009, "type" : "cartoon" }
{ "_id" : ObjectId("59e94bca8eeffc9227fa4a8d"), "imdb" : "tt0171831", "title" :
"Ice Age 4: Continental Drift", "year" : 2012, "type" : "cartoon" }
```



Like FROM clause, specifies the collection to use

Like WHERE clause, specifies which documents to return Projection variables in SELECT clause

How many results to return etc.



The find() Function - projection

- Projection document specifies the fields to return in the documents that match the query filter.
- To return all fields in the matching documents, omit this parameter.
- Form of the projection document:

```
{ field1: <value>, field2: <value> ... }
```

The <value> can be any of the following:

- 1 or true to include the field in the return documents.
- 0 or false to exclude the field.

Projection example:

```
>db.restaurants.find({cuisine: "American", borough: "Bronx"}, {borough:1, restaurant_id:1}).limit(2).sort({restaurant_id:1}).pretty()
```

The _id field is included by default, need to exclude explicitly if you don't want it

```
>db.restaurants.find({cuisine : "American", borough : "Bronx" }, {borough : 1, restaurant_id : 1, _id : 0}).limit(2).sort({restaurant_id : 1}).pretty()
```

```
> db.restaurants.find({cuisine : "American", borough : "Bronx" }, {borough :1, r
estaurant_id : 1, _id : 0}).limit(2).sort({restaurant_id : 1}).pretty()
{ "borough" : "Bronx", "restaurant_id" : "40357217" }
{ "borough" : "Bronx", "restaurant_id" : "40364363" }
>
```



Summary

- JSON document
- MongoDB Document Structure
- Getting start with MongoDB
- MongoDB CRUD Operations
- Data Retrieval in MongoDB